

Chapter 1: Origins, Evolution, and the Cracks in the Foundation

Every developer knows JSON. You've written `{"key": "value"}` thousands of times. You've debugged missing commas, fought with trailing characters, and cursed the lack of comments in configuration files.

But how did we get here? Why does the world's most popular data format have such obvious limitations? And why, despite being "simple," has JSON spawned an entire ecosystem of variants, extensions, and workarounds?

This series explores the JSON you don't know - the one beyond basic syntax. We'll examine binary formats, streaming protocols, validation schemas, RPC layers, and security considerations. But first, we need to understand why JSON exists and where it falls short.

{blurb, class: information} **What XML Had:** Everything built-in (1998-2005)

XML's approach: Monolithic specification with validation (XSD), transformation (XSLT), namespaces, querying (XPath), protocols (SOAP), and security (XML Signature/Encryption) all integrated into one ecosystem.

```
<!-- XML had it all in one place -->
<user xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
      xsi:schemaLocation="user.xsd">
    <name>Alice</name>
    <email>alice@example.com</email>
</user>
```

Benefit: Complete solution with built-in type safety, validation, and extensibility
Cost: Massive complexity, steep learning curve, rigid coupling between features

JSON's approach: Minimal core with separate standards for each need

Architecture shift: Integrated → Modular, Everything built-in → Composable solutions, Monolithic → Ecosystem-driven {/blurb}

The Pre-JSON Dark Ages: XML Everywhere

The Problem Space (Late 1990s)

The web was growing explosively. Websites evolved from static HTML to dynamic applications. Services needed to communicate across networks, applications needed configuration files, and developers needed a way to move structured data between systems.

The requirements were clear: - Human-readable (developers must debug it)
- Machine-parseable (computers must process it) - Language-agnostic (works in

- any programming language) - Supports nested structures (real data has hierarchy)
- Self-describing (data carries its own schema)

XML: The Heavyweight Champion

XML (eXtensible Markup Language) emerged as the answer. By the early 2000s, it dominated:

- XML everywhere: - Configuration files (web.xml, applicationContext.xml)
- SOAP web services (the enterprise standard) - Data exchange (RSS, Atom feeds)
- Document formats (DOCX, SVG) - Build systems (Maven pom.xml, Ant build.xml)

A simple person record in XML:

```
<?xml version="1.0" encoding="UTF-8"?>
<person>
    <name>Alice Johnson</name>
    <email>alice@example.com</email>
    <age>30</age>
    <active>true</active>
    <hobbies>
        <hobby>reading</hobby>
        <hobby>cycling</hobby>
    </hobbies>
</person>
```

Size: 247 bytes

The Road to XML: SGML's Legacy

Before XML dominated, the data format landscape was fragmented and painful. Understanding XML's rise requires understanding what came before.

SGML (Standard Generalized Markup Language, 1986): - Incredibly powerful document markup language - Used for technical documentation, aerospace, defense - So complex it required dedicated specialists - Parsers were proprietary, expensive, and platform-specific - Learning curve measured in months

Example SGML complexity:

```
<!DOCTYPE book [
    <!ELEMENT book - - (title, author+, chapter+)>
    <!ELEMENT title - 0 (#PCDATA)>
    <!ELEMENT author - 0 (#PCDATA)>
    <!ATTLIST author
        role (primary|contributor) #REQUIRED>
]>
<book>
```

```
<title>Data Formats</title>
<author role="primary">Alice Johnson</author>
</book>
```

SGML worked for specialized domains (HTML is actually an SGML application) but was far too complex for general-purpose data exchange.

Early alternatives (1990s): - **CSV:** Great for tabular data, useless for hierarchy - **INI files:** Simple config format, no nesting - **Custom formats:** Every application invented its own (nightmare for interoperability) - **Binary protocols:** Fast but opaque, debugging impossible

The W3C standardization effort (1996-1998):

In 1996, the World Wide Web Consortium (W3C) convened a working group to create “SGML for the web” - a simplified markup language that could: - Retain SGML’s power for structured documents - Be simple enough for average programmers - Work across the internet without specialized tools - Support both documents and data

Key players: - Tim Bray (co-editor, XML specification) - Jean Paoli (Microsoft) - Michael Sperberg-McQueen (University of Illinois) - Jon Bosak (Sun Microsystems, working group chair)

XML 1.0 released: February 10, 1998

The specification aimed for the sweet spot: simpler than SGML, more powerful than HTML, suitable for data and documents.

Early adoption drivers (1998-2000):

Microsoft embraced XML aggressively: - XML support in IE 5 (1999) - XML as configuration format in .NET (2000) - Office XML formats (precursor to DOCX) - MSXML parser bundled with Windows

Sun Microsystems pushed XML for Java: - JAXP (Java API for XML Processing) - XML-based J2EE deployment descriptors - JAXB (Java Architecture for XML Binding)

IBM promoted XML for enterprise: - WebSphere XML tooling - XML database products - Industry consortium participation

The SOAP explosion (2000-2003):

XML’s killer app (and eventual albatross) was SOAP - Simple Object Access Protocol. Microsoft and IBM heavily promoted SOAP as the future of distributed computing, positioning it as the successor to CORBA and DCOM.

SOAP promised: - Language-neutral RPC - HTTP transport (firewall-friendly) - Automatic code generation from WSDL - Enterprise-grade features (security, transactions, reliability)

SOAP delivered: - Massive XML payloads for simple function calls - Complex tooling requirements (WSDL, code generators) - 50+ WS-* specifications (WS-Security, WS-ReliableMessaging, etc.) - Interoperability nightmares despite “standard” protocols

Example SOAP call to add two numbers:

```
<?xml version="1.0"?>
<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
  <soap:Header>
    <wsse:Security xmlns:wsse="http://schemas.xmlsoap.org/ws/2003/06/secext">
      <wsse:UsernameToken>
        <wsse:Username>alice</wsse:Username>
        <wsse:Password Type="PasswordDigest">...</wsse:Password>
      </wsse:UsernameToken>
    </wsse:Security>
  </soap:Header>
  <soap:Body>
    <m:Add xmlns:m="http://example.com/math">
      <m:a>5</m:a>
      <m:b>3</m:b>
    </m:Add>
  </soap:Body>
</soap:Envelope>
```

Size: 500+ bytes to add two numbers. The backlash was inevitable.

Why enterprises loved XML (2000-2005): - Complete specification (nothing undefined) - Schema validation (XSD) enforced contracts - Tool vendors sold expensive IDE plugins - Consultants billed months for SOAP integration - “Enterprise-grade” meant complex, and XML delivered

Why developers hated XML (2003-2010): - 10:1 ratio of markup to data - Namespace hell (xmlns confusion) - XSD schemas more complex than the code - SOAP toolkits measured in megabytes - Debugging massive XML payloads impossible - Editing XML configs required understanding the entire spec

By 2005, the backlash was in full swing. Developers joked about “angle bracket tax” and “XML trauma.” The stage was set for something simpler.

XML’s Strengths

Despite the complexity, XML wasn’t chosen arbitrarily. It had real advantages:

- + **Schema validation** (XSD, DTD, RelaxNG)
- + **Namespaces** (avoid naming conflicts)
- + **XPath** (query language)
- + **XSLT** (transformation)
- + **Comments** (documentation support)

- + **Attributes and elements** (flexible modeling)
- + **Mature tooling** (parsers in every language)

XML's Fatal Flaws: The Monolithic Architecture

But XML's complexity became its downfall. The problem wasn't any single feature - it was the **architectural decision to build everything into one specification**.

XML wasn't just a data format. It was an entire technology stack:

Core XML (parsing and structure): - DOM (Document Object Model) - load entire document into memory - SAX (Simple API for XML) - event-driven streaming parser - StAX (Streaming API for XML) - pull parser - Namespace handling (xmlns declarations) - Entity resolution (external references) - CDATA sections (unparsed character data)

Validation layer (built-in): - DTD (Document Type Definition) - original schema language - XSD (XML Schema Definition) - complex type system - RelaxNG - alternative schema language - Schematron - rule-based validation

Query layer (built-in): - XPath - query language for selecting nodes - XQuery - SQL-like language for XML - XSLT - transformation and templating

Protocol layer (built-in): - SOAP (Simple Object Access Protocol) - WSDL (Web Services Description Language) - WS-Security, WS-ReliableMessaging, WS-AtomicTransaction - 50+ WS-* specifications

The architectural problem: Every XML parser had to support this entire stack. You couldn't use XML without dealing with namespaces. You couldn't validate without learning XSD. You couldn't query without XPath.

The result: - XML parsers: 50,000+ lines of code - XSD validators: Complex type systems rivaling programming languages - SOAP toolkits: Megabytes of libraries just to call a remote function - Learning curve: Months to master the ecosystem

Specific pain points:

Verbosity:

```
<user>
  <name>Alice</name>
  <email>alice@example.com</email>
</user>
```

vs

```
{"name": "Alice", "email": "alice@example.com"}
```

Namespace confusion:

```

<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns:xsd="http://www.w3.org/2001/XMLSchema">
    <soap:Body>
        < GetUser xmlns="http://example.com/users">
            <UserId>123</UserId>
        </ GetUser>
    </soap:Body>
</soap:Envelope>

```

Schema complexity (XSD):

```

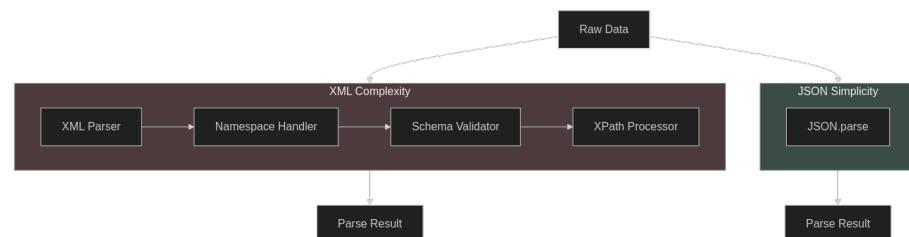
<xs:schema xmlns:xs="http://www.w3.org/2001/XMLSchema">
    <xs:element name="user">
        <xs:complexType>
            <xs:sequence>
                <xs:element name="name" type="xs:string"/>
                <xs:element name="email" type="xs:string"/>
            </xs:sequence>
        </xs:complexType>
    </xs:element>
</xs:schema>

```

vs JSON Schema:

```
{
  "type": "object",
  "properties": {
    "name": {"type": "string"},
    "email": {"type": "string"}
  }
}
```

The real killer: Developer experience. Writing XML by hand was tedious. Reading XML logs was painful. Debugging SOAP requests required specialized tools. The monolithic architecture meant you couldn't use just the parts you needed - it was all or nothing.



{width: 85%}

JSON's Accidental Discovery

Douglas Crockford's Realization (2001)

JSON wasn't invented - it was discovered. Douglas Crockford realized that JavaScript's object literal notation was already a perfect data format:

```
// JavaScript code that's also data
var person = {
    name: "Alice Johnson",
    email: "alice@example.com",
    age: 30,
    active: true,
    hobbies: ["reading", "cycling"]
};
```

Key insight: This notation was:
- Already in JavaScript engines (browsers everywhere)
- Minimal syntax (no closing tags)
- Easy to parse (recursive descent parser is ~500 lines)
- Human-readable
- Machine-friendly

The Same Data in JSON

```
{
    "name": "Alice Johnson",
    "email": "alice@example.com",
    "age": 30,
    "active": true,
    "hobbies": ["reading", "cycling"]
}
```

Size: 129 bytes (52% smaller than XML)

The Simplicity Revolution

JSON's radical simplification:

Six data types: 1. object - { "key": "value" } 2. array - [1, 2, 3] 3. string - "text" 4. number - 123 or 123.45 5. boolean - true or false 6. null - null

That's it. No attributes. No namespaces. No CDATA sections. No processing instructions.

Browser Native Support

The killer feature:

```
// Parse JSON (browsers built-in)
var data = JSON.parse(jsonString);
```

```
// Generate JSON
var json = JSON.stringify(data);
```

No XML parser library needed. No SAX vs DOM decision. Just two functions.

Evolution of Data Formats:

Year	Milestone	Significance
1998	XML 1.0 Specification	Monolithic data format dominates
	SOAP begins development	Enterprise web services standard
2001	JSON discovered by Crockford	JavaScript object notation as data format
	First JSON parsers appear	Language support begins
	JSON used in AJAX applications	Google Maps, Web 2.0 movement
2006	RFC 4627 - JSON specification	JSON becomes formal standard
2013	RFC 7159 - Updated JSON spec	Improved specification
2017	ECMA-404 standard	Dual standardization
	RFC 8259 - Current JSON standard	Stable, mature specification
2020+	JSON dominates REST APIs	Clear winner for data interchange
	JSON Schema, JSONB, JSONL	Complete modular ecosystem
	JSON ecosystem mature	Production-ready tooling

Why JSON Won

1. The AJAX Revolution (2005)

February 8, 2005: Google launches Google Maps. The web would never be the same.

Before Google Maps, web pages were static. Click a link, wait for full page reload, repeat. Google Maps let you drag the map, and tiles loaded dynamically without page refresh. The effect was magical - it felt like a desktop application running in a browser.

The technical breakthrough: XMLHttpRequest

Google Maps (and Gmail before it) used XMLHttpRequest - a Microsoft invention from 1999 that had languished unused. Jesse James Garrett coined the term

“AJAX” (Asynchronous JavaScript and XML) in his February 2005 essay.

The immediate problem: XML was terrible for this use case.

Why XML failed for AJAX:

1. Parsing performance

```
// XML parsing in browser (2005)
const xhr = new XMLHttpRequest();
xhr.open('GET', '/api/markers');
xhr.onload = function() {
    const xml = xhr.responseXML; // Parse entire document
    const markers = xml.getElementsByTagName('marker');
    // Traverse DOM tree
    for (let i = 0; i < markers.length; i++) {
        const lat = markers[i].getAttribute('lat');
        const lng = markers[i].getAttribute('lng');
        // Process marker...
    }
};
```

Parsing time for 1000 markers: ~200ms (IE 6), ~150ms (Firefox 1.5)

```
// JSON parsing in browser (2005)
const xhr = new XMLHttpRequest();
xhr.open('GET', '/api/markers');
xhr.onload = function() {
    const markers = JSON.parse(xhr.responseText); // Single operation
    // Directly access array
    markers.forEach(marker => {
        const {lat, lng} = marker;
        // Process marker...
    });
};
```

Parsing time for 1000 markers: ~15ms (IE 6), ~8ms (Firefox 1.5)

JSON was 10-20x faster because `JSON.parse()` was implemented in native C code while XML parsing required building an entire DOM tree in JavaScript.

2. Bandwidth costs (critical in 2005)

Average connection speeds in 2005: - Dial-up: 56 kbps (still 35% of users) - Broadband: 384 kbps - 1 Mbps - Mobile: 64-144 kbps (EDGE)

Example: 100 map markers

XML response:

```
<?xml version="1.0" encoding="UTF-8"?>
<markers>
```

```

<marker lat="37.7749" lng="-122.4194" name="San Francisco"/>
<marker lat="34.0522" lng="-118.2437" name="Los Angeles"/>
<!-- ... 98 more markers ... -->
</markers>

```

Size: ~8,500 bytes (with full tags)

JSON response:

```

{
  "markers": [
    {"lat": 37.7749, "lng": -122.4194, "name": "San Francisco"},
    {"lat": 34.0522, "lng": -118.2437, "name": "Los Angeles"}
  ]
}

```

Size: ~3,200 bytes (62% smaller)

At 56 kbps (dial-up): - XML: 1.2 seconds to transfer - JSON: 0.45 seconds to transfer

The difference was user-perceptible. Google Maps with JSON felt snappy. With XML, it would have felt sluggish.

3. Developer ergonomics

XML:

```

// Extract data from XML (verbose, error-prone)
const name = marker.getElementsByTagName('name')[0].textContent;
const lat = parseFloat(marker.getAttribute('lat'));

```

JSON:

```

// Extract data from JSON (natural JavaScript)
const {name, lat} = marker;

```

The irony of “AJAX”

Despite the name (“Asynchronous JavaScript and XML”), JSON became the dominant data format for AJAX applications by 2006-2007. Developers joked about “AJAJ” (Asynchronous JavaScript and JSON) but the AJAX name stuck.

XMLHttpRequest naming irony: The API kept the XML name even though 90%+ of modern usage is for JSON. This historical artifact confuses new developers to this day.

Real-world impact:

By 2007: - Google Maps API officially supported JSON responses - Yahoo APIs defaulted to JSON - Flickr API offered JSON alongside XML - Twitter API (launched 2006) was JSON-first

The transition was swift because JSON solved real pain points:

- 60%+ bandwidth reduction
- 10x+ parsing performance improvement
- Native JavaScript integration
- No XML parser library required

The turning point: When Google - the largest web company - chose JSON over XML for their flagship AJAX applications, the industry followed. XML remained for documents and configuration, but for APIs carrying data, JSON won decisively.

2. REST vs SOAP

REST APIs adopted JSON as the default format:

SOAP request (XML):

```
<?xml version="1.0"?>
<soap:Envelope xmlns:soap="http://www.w3.org/2003/05/soap-envelope">
  <soap:Header>
  </soap:Header>
  <soap:Body>
    <m:GetUser xmlns:m="http://example.com/users">
      <m:UserId>123</m:UserId>
    </m:GetUser>
  </soap:Body>
</soap:Envelope>
```

REST request (JSON):

```
GET /users/123
Accept: application/json
```

REST response:

```
{
  "id": 123,
  "name": "Alice Johnson",
  "email": "alice@example.com"
}
```

The difference was stark. REST + JSON became the de facto standard for web APIs.

3. NoSQL Movement (2009+)

MongoDB, CouchDB, and other NoSQL databases chose JSON-like formats:

```
// MongoDB document (BSON internally)
{
  "_id": ObjectId("507f1f77bcf86cd799439011"),
  "name": "Alice Johnson",
  "email": "alice@example.com",
```

```
    "created": ISODate("2023-01-15T10:30:00Z")
}
```

Why JSON for databases: - Schema flexibility (add fields without migrations)
- Direct JavaScript integration - Document model matches JSON structure
- Query results are already in API format

4. Configuration Files

JSON displaced XML in configuration:

package.json (Node.js):

```
{
  "name": "my-app",
  "version": "1.0.0",
  "dependencies": {
    "express": "^4.18.0"
  }
}
```

tsconfig.json (TypeScript):

```
{
  "compilerOptions": {
    "target": "ES2020",
    "module": "commonjs",
    "strict": true
  }
}
```

Developers preferred JSON over XML for configuration because it was easier to read and edit.

5. The Complete Comparison: XML vs JSON in Practice

To understand why JSON displaced XML so thoroughly, let's compare them across every dimension that matters for real-world development.

The same data in both formats:

XML:

```
<?xml version="1.0" encoding="UTF-8"?>
<user xmlns="http://example.com/user" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
  <id>user-5f9d88c</id>
  <username>alice</username>
  <email>alice@example.com</email>
  <profile>
    <firstName>Alice</firstName>
    <lastName>Johnson</lastName>
```

```

<age>30</age>
</profile>
<verified>true</verified>
<tags>
  <tag>golang</tag>
  <tag>rust</tag>
  <tag>distributed-systems</tag>
</tags>
<metadata>
  <created>2023-01-15T10:30:00Z</created>
  <lastLogin>2023-12-01T09:15:00Z</lastLogin>
</metadata>
</user>

```

Size: 512 bytes

JSON:

```
{
  "id": "user-5f9d88c",
  "username": "alice",
  "email": "alice@example.com",
  "profile": {
    "firstName": "Alice",
    "lastName": "Johnson",
    "age": 30
  },
  "verified": true,
  "tags": ["golang", "rust", "distributed-systems"],
  "metadata": {
    "created": "2023-01-15T10:30:00Z",
    "lastLogin": "2023-12-01T09:15:00Z"
  }
}
```

Size: 312 bytes (39% smaller)

Parsing complexity comparison:

XML parsing (JavaScript):

```
// Requires XML parser and DOM traversal
const parser = new DOMParser();
const doc = parser.parseFromString(xmlString, 'text/xml');

// Extract nested data requires multiple steps
const id = doc.getElementsByTagName('id')[0].textContent;
const firstName = doc.querySelector('profile firstName').textContent;
const tags = Array.from(doc.querySelectorAll('tags tag'))
```

```
.map(tag => tag.textContent);
const age = parseInt(doc.querySelector('profile age').textContent);

// Type conversion manual
const verified = doc.getElementsByTagName('verified')[0].textContent === 'true';
```

Lines of code: 10

Parsing time (1000 objects): ~180ms

Memory allocation: DOM tree + JavaScript objects

JSON parsing (JavaScript):

```
// Native operation, single line
const user = JSON.parse(jsonString);

// Direct property access
const {id, profile: {firstName}, tags, age, verified} = user;
```

Lines of code: 2

Parsing time (1000 objects): ~12ms (15x faster)

Memory allocation: JavaScript objects only

Ecosystem size comparison (as of 2010):

XML ecosystem (built-in specifications): - Core: XML 1.0, XML Namespaces, XML Base - Validation: DTD, XSD, RelaxNG, Schematron (4 competing standards) - Transformation: XSLT 1.0, XSLT 2.0, XSLT 3.0 - Querying: XPath 1.0, XPath 2.0, XQuery - Protocols: SOAP 1.1, SOAP 1.2, 50+ WS-* specs - Security: XML Signature, XML Encryption - Total: 100+ W3C specifications

JSON ecosystem (2010): - Core: RFC 4627 (JSON specification) - Validation: JSON Schema (emerging) - Total: 1 specification + 1 emerging standard

Learning curve measurement (time to productivity):

XML mastery path: - Week 1: Basic XML syntax, well-formedness - Week 2: Namespaces, attributes vs elements - Week 3: XSD schema basics - Week 4: XPath querying - Month 2: XSLT transformation - Month 3: SOAP web services - Month 4+: WS-Security, advanced patterns

Total time to proficiency: 3-6 months

JSON mastery path: - Day 1: Basic syntax, nesting - Day 2: Use with APIs - Week 1: Proficient

Total time to proficiency: 1 week

When XML was actually better:

Despite JSON's advantages, XML genuinely excelled in specific domains:

1. Document markup (XML's original purpose):

```

<book>
  <chapter id="1">
    <title>Introduction</title>
    <paragraph>
      This is <emphasis>important</emphasis> text with
      <footnote id="fn1">Additional context</footnote>.
    </paragraph>
  </chapter>
</book>

```

Mixed content (text + markup) is natural in XML, awkward in JSON. HTML, DocBook, and EPUB use XML for good reason.

2. Namespaces for mixing vocabularies:

```

<html xmlns="http://www.w3.org/1999/xhtml"
      xmlns:svg="http://www.w3.org/2000/svg">
  <body>
    <svg:svg width="100" height="100">
      <svg:circle cx="50" cy="50" r="40"/>
    </svg:svg>
  </body>
</html>

```

XML namespaces allow mixing HTML and SVG without ambiguity. JSON has no equivalent.

3. Schema evolution with extensibility:

```

<user xmlns:ext="http://example.com/extensions">
  <name>Alice</name>
  <ext:customField>value</ext:customField>
</user>

```

Old XML parsers ignore unknown namespaces gracefully. JSON parsers must handle unknown fields explicitly.

4. Comments in production data:

```

<config>
  <!-- TODO: Update this before release -->
  <setting>value</setting>
</config>

```

XML's native comment support is genuinely useful for configuration files with documentation. JSON's lack of comments is a real limitation.

The verdict: Context matters

JSON didn't "win" universally - it won for **data interchange** specifically: - APIs carrying structured data: JSON dominates - Configuration files: JSON/YAML (YAML adds comments JSON lacks) - Databases: JSON-based (MongoDB,

PostgreSQL JSONB) - Document markup: XML still standard (EPUB, DocBook, SVG) - Office formats: XML (DOCX, XLSX, PPTX)

JSON succeeded because it **matched the right use case** (data interchange) with the **right architecture** (minimal, composable) at the **right time** (AJAX revolution, REST APIs emerging).

6. Language Support Explosion

By 2010, every major language had JSON support:

Go:

```
import "encoding/json"

type Person struct {
    Name string `json:"name"`
    Email string `json:"email"`
    Age int `json:"age"`
}

json.Marshal(person) // encode
json.Unmarshal(data, &person) // decode
```

Python:

```
import json

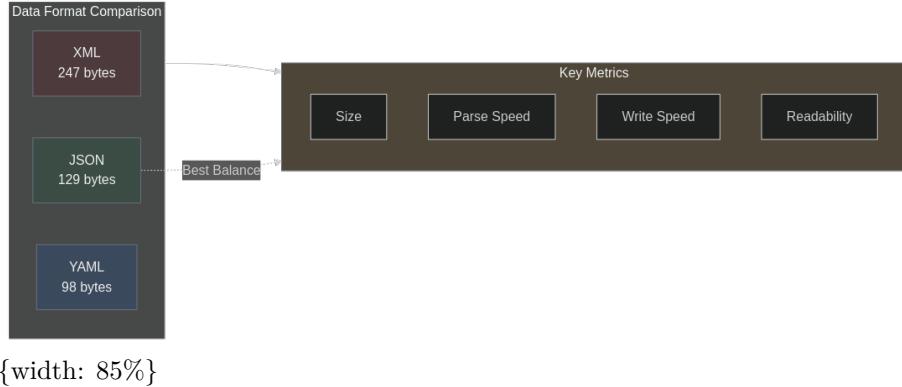
person = {"name": "Alice", "email": "alice@example.com"}
json.dumps(person) # encode
json.loads(data) # decode
```

Java:

```
import com.fasterxml.jackson.databind.ObjectMapper;

ObjectMapper mapper = new ObjectMapper();
String json = mapper.writeValueAsString(person); // encode
Person person = mapper.readValue(json, Person.class); // decode

{blurb, class: information} The Ecosystem Effect: Once every language had JSON support, it became the obvious choice for data interchange. Network effects made JSON the default - not because it was technically superior, but because it was universally supported. {/blurb}
```



JSON's Fundamental Weaknesses

Now we reach the core problem. JSON won because it was simple. But that simplicity came with trade-offs that become painful at scale.

1. No Schema or Validation

The problem:

```
{
  "name": "Alice",
  "age": "30"
}
```

Is `age` a string or a number? Both are valid JSON. The parser accepts both. Your application crashes when it expects a number.

Real-world consequences: - API breaking changes go undetected - Invalid data passes validation - Runtime errors instead of compile-time checks - Documentation is separate from data format - Client-server contract is implicit, not explicit

2. No Date/Time Type

JSON has no standard way to represent dates:

```
{
  "created": "2023-01-15"
}
{
  "created": "2023-01-15T10:30:00Z"
}
```

```
{  
  "created": 1673780400  
}
```

All are valid JSON. Which format do you use? ISO 8601 string? Unix timestamp?
Custom format?

Every project reinvents this. Libraries make assumptions. APIs document their chosen format. Parsing errors happen when formats don't match.

3. Number Precision Issues

JavaScript uses IEEE 754 double-precision floats for all numbers:

```
// JavaScript  
console.log(9007199254740992 + 1); // 9007199254740992  
// Lost precision!
```

{blurb, class: error} **Critical Production Issue:** JSON's number type causes real-world failures:
- Database IDs beyond 2^{53} silently corrupt (Snowflake IDs, Twitter IDs)
- Financial calculations lose cents (\$1234.56 becomes \$1234.5599999999)
- Timestamps break (millisecond precision lost after 2^{53})
- Different languages parse differently (Python preserves precision, JavaScript doesn't)

This isn't theoretical - major APIs (Twitter, Stripe, GitHub) return large IDs as strings to prevent JavaScript corruption. If your API has >10M records with auto-increment IDs, you WILL hit this. {/blurb}

Problems: - Large integers lose precision (database IDs, timestamps)
- No distinction between integer and float
- Different languages handle this differently
- Financial calculations require special handling

Common workaround:

```
{  
  "id": "9007199254740993",  
  "balance": "1234.56"  
}
```

Represent numbers as strings to preserve precision. But now you need custom parsing logic.

Real-world examples:

```
// Twitter API returns IDs as strings  
{  
  "id": 1234567890123456789,      // Unsafe in JavaScript  
  "id_str": "1234567890123456789" // Always use this  
}
```

```

// Stripe amounts are integers (cents)
{
  "amount": 123456, // \$1234.56 as integer cents
  "currency": "usd"
}

// Shopify order numbers as strings
{
  "order_number": "1001", // String to avoid precision issues
  "total": "29.99" // String for exact decimal
}

```

4. No Comments

You cannot add comments to JSON:

```
{
  "port": 8080,
  "debug": true
}
```

Why is debug enabled? What does this configuration do? You can't document it in the file itself.

Workarounds:

```
{
  "_comment": "Enable debug mode in development",
  "debug": true
}
```

Use fake fields for comments. But parsers still process these as data.

5. No Binary Data Support

JSON is text-based. Binary data must be encoded:

```
{
  "image": "iVBORw0KGgoAAAANSUhEUgAAAAEAAAABCAYAAQfFcSJAAADU1EQVR42mNk+M9QDwADhgGAWjR9awAA"
}
```

Problems: - Base64 encoding increases size by ~33% - Additional encoding/decoding overhead - Not efficient for large binary files

6. Verbose for Large Datasets

Repeated field names add significant overhead:

```
[
  {"id": 1, "name": "Alice", "email": "alice@example.com"},
```

```

  {"id": 2, "name": "Bob", "email": "bob@example.com"},  

  {"id": 3, "name": "Carol", "email": "carol@example.com"}  

]

```

Field names (“id”, “name”, “email”) repeat for every record. In a 100,000 row dataset, this is wasteful.

CSV alternative (for comparison):

```

id,name,email  

1,Alice,alice@example.com  

2,Bob,bob@example.com  

3,Carol,carol@example.com

```

More compact, but loses type information and nested structure support.

7. No Circular References

JSON cannot represent circular references:

```

// JavaScript object
let person = {name: "Alice"};
let company = {name: "Acme Corp", ceo: person};
person.employer = company; // Circular reference

JSON.stringify(person);
// TypeError: Converting circular structure to JSON

```

You must manually break cycles or use a serialization library that detects and handles them.

{blurb, class: warning} **Critical Insight:** JSON’s weaknesses aren’t bugs - they’re consequences of extreme simplification. Every missing feature (schemas, comments, binary support) was left out intentionally to keep the format minimal.
{/blurb}

The Format Comparison Landscape

Let’s compare JSON to its alternatives across key dimensions:

Feature	JSON	XML	YAML	TOML	Protocol Buffers
Human-readable	Yes	Yes	Yes	Yes	No
Schema validation	No*	Yes	No	No	Yes
Comments	No	Yes	Yes	Yes	No

Feature	JSON	XML	YAML	TOML	Protocol Buffers
Binary support	No	No	No	No	Yes
Date types	No	No	No	Yes	Yes
Size efficiency	Medium	Large	Medium	Medium	Small
Parse speed	Fast	Slow	Medium	Medium	Very Fast
Language support	Universal	Universal	Wide	Growing	Wide
Nested structures	Yes	Yes	Yes	Limited	Yes
Trailing commas	No	N/A	Yes	Yes	N/A
Type safety	No	Yes	No	Partial	Yes

*JSON Schema provides validation but isn't part of JSON itself.



{width: 85%}

When NOT to Use JSON

Despite JSON's dominance, there are clear cases where alternatives are better:

1. High-Performance Systems → Protocol Buffers, FlatBuffers

When you're handling millions of requests per second, Protocol Buffers offer compelling advantages:

```
message Person {
    string name = 1;
    string email = 2;
```

```
    int32 age = 3;  
}
```

Benefits: - 3-10x smaller than JSON - 5-20x faster to parse - Schema enforced at compile time - Backward/forward compatibility built-in

Trade-off: Not human-readable, requires schema compilation.

Read more: Understanding Protocol Buffers: Part 1

2. Human-Edited Configuration → YAML, TOML, JSON5

When developers edit config files frequently:

TOML:

```
[database]  
host = "localhost"  
port = 5432  
  
# Connection pool settings  
[database.pool]  
max_connections = 100  
min_connections = 10
```

YAML:

```
database:  
  host: localhost  
  port: 5432  
  pool:  
    max_connections: 100  
    min_connections: 10 # Minimum pool size
```

Benefits: Comments, less syntax noise, more readable.

Trade-off: YAML has subtle parsing gotchas (indentation, special values like no/yes).

3. Large Tabular Datasets → CSV, Parquet, Arrow

For analytics and data pipelines:

```
id,name,email,created  
1,Alice,alice@example.com,2023-01-15  
2,Bob,bob@example.com,2023-01-16
```

Benefits: Much more compact, streaming-friendly, tooling optimized for analysis.

Trade-off: No nested structures, limited type information.

4. Document Storage → BSON, MessagePack

When JSON-like flexibility meets binary efficiency:

```
// MongoDB (BSON)
{
  _id: ObjectId("507f1f77bcf86cd799439011"),
  name: "Alice",
  created: ISODate("2023-01-15T10:30:00Z"),
  avatar: BinData(0, "base64data...")
}
```

Benefits: Native date types, binary data support, efficient storage.

Trade-off: Binary format, language-specific implementations.

The Evolution: JSON's Ecosystem Response

JSON's limitations didn't kill it. Instead, an entire ecosystem evolved to address the weaknesses while preserving the core simplicity:

{blurb, class: warning} **The Architectural Choice:** XML's completeness was a weakness - validation, namespaces, transformation, and querying were built into one monolithic specification. **Every XML parser needed to support everything, making the system rigid and complex.** JSON chose the opposite path: radical incompleteness. The core format has no validation, no binary support, no streaming, no protocol conventions. Each gap would be filled by modular, composable solutions that could evolve independently. {/blurb}

1. Validation Layer: JSON Schema

Problem: No built-in validation

Solution: External schema language

```
{
  "type": "object",
  "properties": {
    "name": {"type": "string", "minLength": 1},
    "email": {"type": "string", "format": "email"},
    "age": {"type": "integer", "minimum": 0}
  },
  "required": ["name", "email"]
}
```

{blurb, class: tip} **Transformation:** This single innovation transformed JSON from “hope the data is correct” to “validate at runtime with strict schemas.” JSON Schema adds the type safety layer that JSON itself deliberately omitted. {/blurb}

Next article: Part 2 dives deep into JSON Schema - how it works, why it matters, and how it solves JSON's validation problem.

2. Binary Variants: JSONB, BSON, MessagePack

Problem: Text format is inefficient

Solution: Binary encoding with JSON-like structure

These formats maintain JSON's structure while using efficient binary serialization:

- **PostgreSQL JSONB:** Decomposed binary format, indexable, faster queries
- **MongoDB BSON:** Binary JSON with extended types
- **MessagePack:** Universal binary serialization

3. Streaming Format: JSON Lines (JSONL)

Problem: JSON arrays don't stream

Solution: Newline-delimited JSON objects

```
{"id": 1, "name": "Alice"}  
{"id": 2, "name": "Bob"}  
{"id": 3, "name": "Carol"}
```

Each line is independent, enabling streaming, log files, and Unix pipeline processing.

4. Protocol Layer: JSON-RPC

Problem: No standard RPC convention

Solution: Structured request/response format

```
{  
  "jsonrpc": "2.0",  
  "method": "getUser",  
  "params": {"id": 123},  
  "id": 1  
}
```

Used by Ethereum, LSP (Language Server Protocol), and many other systems.

5. Human-Friendly Variants: JSON5, HJSON

Problem: No comments, strict syntax

Solution: Relaxed JSON with comments and trailing commas

JSON deliberately omitted developer-friendly features to stay minimal. For machine-to-machine communication, this is fine. For configuration files humans edit daily, it's painful.

JSON5 adds convenience features while maintaining JSON compatibility:

```
{
  // Single-line comments
  /* Multi-line comments */

  name: 'my-app',           // Unquoted keys
  port: 8080,
  features: {
    debug: false,          // Trailing commas OK
    maxConnections: 1_000,   // Numeric separators
  },
  description: 'JSON5 supports cleaner syntax',
}
```

HJSON goes further with extreme readability:

```
{
  # Hash comments (like YAML)

  # Quotes optional for strings
  name: my-app
  port: 8080

  # Commas optional
  features: {
    debug: false
    maxConnections: 1000
  }

  # Multi-line strings without escaping
  description:
  '''
  This is a naturally
  multi-line description
  '''
}
```

Comparison:

Feature	JSON	JSON5	HJSON	YAML	TOML
Comments	No	Yes	Yes	Yes	Yes
Trailing commas	No	Yes	Yes	N/A	N/A
Unquoted keys	No	Yes	Yes	Yes	Yes
Unquoted strings	No	No	Yes	Yes	Yes
Native browser support	Yes	No	No	No	No
Designed for configs	No	Partial	Yes	Yes	Yes

When to use:

JSON5: - VSCode settings (`.vscode/settings.json5`) - Build tool configs where JSON is expected - Need comments but want JSON compatibility

HJSON: - Developer-facing configs prioritizing readability - Local development settings - Documentation examples

Standard JSON: - APIs and data interchange - Production configs (parsed by machines) - Anything needing browser/native support

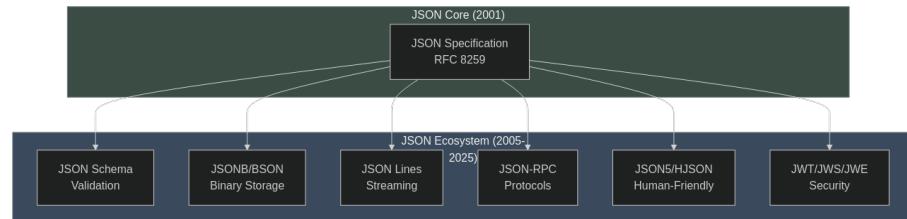
Why they're niche: Unlike JSON Schema (essential for validation) or JSONB (essential for performance), JSON5/HJSON solve a convenience problem that YAML and TOML also solve. Most teams choose YAML or TOML for configuration files - they were designed for this purpose from the start and have broader ecosystem support.

{blurb, class: information} **The Configuration Choice:** For human-edited configs, the ecosystem offers multiple solutions - JSON5, HJSON, YAML, TOML. Each makes different trade-offs between readability, features, and compatibility. JSON5 stays closest to JSON, YAML is most popular, TOML is clearest for nested config. The choice depends on your team's preferences and tooling.
{/blurb}

6. Security Layer: JWS, JWE

Problem: No built-in security

Solution: JSON Web Signatures and Encryption standards



{width: 85%}

Running Example: Building a User API

Throughout this series, we'll follow a single use case: **a User API for a social platform**. Each part will show how that layer of the ecosystem improves this real-world scenario.

The scenario: - REST API for user management - 10 million users in PostgreSQL - Mobile and web clients - Need authentication, validation, performance, and security

Part 1 (this article): The basic JSON structure

```
{  
  "id": "user-5f9d88c",  
  "username": "alice",  
  "email": "alice@example.com",  
  "created": "2023-01-15T10:30:00Z",  
  "bio": "Software engineer",  
  "followers": 1234,  
  "verified": true  
}
```

What's missing: - No validation (what if email is invalid?) - Inefficient storage (text format repeated 10M times) - Can't stream user exports (arrays don't stream) - No authentication (how do we secure this?) - No protocol (how do clients call getUserId?)

The journey ahead: - **Part 2:** Add JSON Schema validation for type safety - **Part 3:** Store users in PostgreSQL JSONB for performance - **Part 5:** Add JSON-RPC protocol for structured API calls - **Part 5:** Export users with JSON Lines for streaming - **Part 6:** Secure API with JWT authentication

This single API will demonstrate how each ecosystem layer solves a real problem.

Conclusion: JSON's Success Through Simplicity

JSON won not because it was perfect, but because it was simple enough to understand, implement, and adopt universally. Its weaknesses are real, but they're addressable through layered solutions.

What made JSON win: - Minimal syntax (6 data types, simple rules) - Browser native support (JSON.parse/stringify) - Perfect timing (AJAX era, REST movement) - Universal language support (parsers in everything) - Good enough for most use cases

What JSON lacks: - Schema validation (solved by JSON Schema) - Binary efficiency (solved by JSONB, BSON, MessagePack) - Streaming support (solved by JSON Lines) - Protocol conventions (solved by JSON-RPC) - Human-friendly syntax (solved by JSON5, HJSON)

The JSON ecosystem evolved to patch these gaps while preserving the core simplicity that made JSON successful.

{blurb, class: information} **Series Roadmap:** This series explores the JSON ecosystem: - **Part 1** (this article): Origins and fundamental weaknesses - **Part 2:** JSON Schema - validation, types, and contracts - **Part 3:** Binary JSON formats - JSONB, BSON, MessagePack - **Part 6:** Streaming JSON - JSON

Lines and large datasets - **Part 5:** JSON-RPC and protocol layers - **Part 6:** Security - JWT, canonicalization, and attacks {/blurb}

In Part 2, we'll solve JSON's most critical weakness: the lack of validation. JSON Schema transforms JSON from "untyped text" into "strongly validated contracts" without sacrificing simplicity. We'll explore how to define schemas, validate data at runtime, generate code from schemas, and integrate validation into your entire stack.

The core problem JSON Schema solves: How do you maintain the simplicity of JSON while gaining the safety of typed, validated data?

Next: You Don't Know JSON: Part 2 - JSON Schema and the Art of Validation

Further Reading

Specifications: - RFC 8259 - JSON Standard - ECMA-404 - JSON Data Interchange Format

Historical: - Douglas Crockford - The JSON Saga - JSON.org - Introducing JSON

Comparisons: - XML vs JSON Performance Benchmarks - Protocol Buffers vs JSON